

CMPUT201W20B2 Week 13

Abram Hindle

April 6, 2020

Contents

1	Week13	1
1.1	Copyright Statement	1
1.1.1	License	2
1.1.2	Hazel Code is licensed under AGPL3.0+	2
1.2	Alternative version	2
1.3	Init ORG-MODE	2
1.3.1	Org export	3
1.3.2	Org Template	3
1.4	Remember how to compile?	3
1.5	IO	3
1.5.1	Streams	3
1.5.2	Files	8
1.5.3	Command line arguments	17
1.5.4	mmap()	18
1.6	References	20

1 Week13

<https://github.com/abramhindle/CMPUT201W20B2-public/tree/master/week12>

1.1 Copyright Statement

If you are in CMPUT201 at UAlberta this code is released in the public domain to you.

Otherwise it is (c) 2020 Abram Hindle, Hazel Campbell AGPL3.0+

1.1.1 License

CMPUT 201 C Notes Copyright (C) 2020 Abram Hindle, Hazel Campbell

This program is free software: you can redistribute it and/or modify it under the terms of the GNU Affero General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Affero General Public License for more details.

You should have received a copy of the GNU Affero General Public License along with this program. If not, see <https://www.gnu.org/licenses/>.

1.1.2 Hazel Code is licensed under AGPL3.0+

Hazel's code is also found here <https://github.com/hazelybell/examples/tree/C-2020-01>

Hazel code is licensed: The example code is licensed under the AGPL3+ license, unless otherwise noted.

1.2 Alternative version

Checkout the .txt, the .pdf, and the .html version

1.3 Init ORG-MODE

```
;; I need this for org-mode to work well
;; If we have a new org-mode use ob-shell
;; otherwise use ob-sh --- but not both!
(if (require 'ob-shell nil 'noerror)
    (progn
      (org-babel-do-load-languages 'org-babel-load-languages '((shell . t))))
    (progn
      (require 'ob-sh)
      (org-babel-do-load-languages 'org-babel-load-languages '((sh . t))))
  (org-babel-do-load-languages 'org-babel-load-languages '((C . t)))
  (org-babel-do-load-languages 'org-babel-load-languages '((python . t)))
  (setq org-src-fontify-natively t)
  (setq org-confirm-babel-evaluate nil) ;; danger!
  (custom-set-faces
```

```
'(org-block ((t (:inherit shadow :foreground "black"))))
'(org-code ((t (:inherit shadow :foreground "black")))))
```

1.3.1 Org export

```
(org-html-export-to-html)
(org-latex-export-to-pdf)
(org-ascii-export-to-ascii)
```

1.3.2 Org Template

Copy and paste this to demo C

```
#include <stdio.h>

int main(int argc, char**argv) {
    return 0;
}
```

1.4 Remember how to compile?

```
gcc -std=c99 -pedantic -Wall -Wextra -ftrapv -ggdb3 -o programname pro-
gramname.c
```

1.5 IO

stdio.h in C contains numerous IO routines.

You use it primarily for printf and scanf.

1.5.1 Streams

Programs that run in the UNIX terminal have 3 main streams:

- stdin
 - standard in or standard input to read
 - shell: ‘<’ ‘|’
 - C: ‘gets, getchar, scanf, fgets(stdin,...), read(stdin,...) , ...’
- stdout
 - standard out or standard output to write out to the terminals

- shell: ‘>’ ‘|’
- C: ‘puts, printf, fputs(stdout,...), fputc(stdout,...), ...’
- stderr:
 - standard err or standard error to write out to terminals but not modify the main output
 - ‘2>’ ‘|&’ or ‘2>&1 |’
 - C: ‘fputs(stderr,...), fputc(stderr,...), fprintf(stderr,...), ...’

They are called streams because you serially output information to them. And multiple sources can write to the stream. It’s like talking or a stream of consciousness. 1 byte after another.

1. shell

Typically a terminal will mix stdout and stderr.

You can type in input to standard input.

You can "pipe" input to standard input: |

You can redirect file input to standard input: <

Example: using a pipe to pipe the string ‘ALL CAPS’ through the tr program to lower case it.

```
echo ALL CAPS | tr '[:upper:]' '[:lower:]'
```

```
all caps
```

tr is a translation program it takes characters from 1 argument and turns them into another.

```
echo ALL CAPS | tr 'ALC' 'ODP'
```

```
ODD POPS
```

We can make files by redirecting stdout to a file

```
echo ALL CAPS > allcaps.txt
cat allcaps.txt | tr 'AL' 'OP'
echo From Ontario, catch those # just echo to stdout
tr 'AC' 'IR' < allcaps.txt
```

OPP COPS
From Ontario, catch those
ILL RIPS

We can filter arbitrary commands:

```
ls | tr '[:lower:]' '[:upper:]'
```

20.TXT
30.TXT
ALLCAPS.TXT
ARGV
ARGV.C
ARGV-NEW
ARGVRAND
ARGVRAND.C
AUTO
BINARY.BIN
BINARYREAD.C
BINARYWRITE.C
COOLBEARS.TXT
FFLUSHRANDR
FFLUSHRANDR.C
FFLUSHREADER
FFLUSHREADER.C
FFLUSH.SH
FFLUSH.TXT
FGETS.TXT
FPRINTF.TXT
PRESENTATION.HTML
PRESENTATION.ORG
PRESENTATION.ORG~
STDIN-EXAMPLE
STDIN-EXAMPLE.C
STDOUT-EXAMPLE
STDOUT-EXAMPLE.C

We can chain pipes:

```

tr 'AC' 'IR' < allcaps.txt | sed -e 's/LL/LK/'
tr 'AC' 'IR' < allcaps.txt | sed -e 's/LL/LK/' | \
    sed -e 's/^/S/'
# we can chain commands together
tr 'AC' 'IR' < allcaps.txt | sed -e 's/LL/LK/' | \
    sed -e 's/^/S/' | \
    sed -e 's/K /K T/'
ls | grep .org | sort

```

```

ILK RIPS
SILK RIPS
SILK TRIPS
presentation.org
presentation.org~

```

(a) stderr & shell

```

ls -l missing
exit 0

```

Where is it?

```

ls -l missing 2>&1
exit 0

```

```

ls: cannot access 'missing': No such file or directory

```

Once we redirect stderr to stdout we can pipe it and manipulate it!

```

ls -l missing 2>&1 | tr '[:lower:]' '[:upper:]'

```

```

LS: CANNOT ACCESS 'MISSING': NO SUCH FILE OR DIRECTORY

```

Or perhaps we don't want to see the error

```

ls -l missing 2> /dev/null
exit 0

```

Maybe we just want stderr

```

ls -l *.org missing 2>&1 > /dev/null
exit 0

```

```
ls: cannot access 'missing': No such file or directory
```

Maybe we just want BOTH

```
ls -l *.org missing 2>&1
exit 0
```

```
ls: cannot access 'missing': No such file or directory
-rw-r--r-- 1 hindle1 hindle1 22703 Apr  6 23:32 presentation.org
```

2. C

(a) output

```
#include <stdio.h>
int main() {
    printf("OK this is to stdout!\n");
    fprintf(stdout,"OK this is to stdout as well!\n");
    fprintf(stderr,"OK this is to stderr!\n");
    return 0;
}
```

```
OK this is to stdout!
OK this is to stdout as well!
```

Hmmm org-mode ignores stderr

```
gcc -std=c99 -pedantic -Wall -Wextra -ftrapv -ggdb3 -o stdout-example stdout-
./stdout-example 2>&1
```

```
OK this is to stderr!
OK this is to stdout!
OK this is to stdout as well!
```

Ah now it appears

(b) input

```
#include <stdio.h>
int main() {
    int input;
    if (scanf("%d", &input)!=1) abort();
    fprintf(stdout, "From stdin %d\n", input);
    fprintf(stderr, "ERR: From stdin %d\n", input);
    return 0;
}
```

```
gcc -std=c99 -pedantic -Wall -Wextra -ftrapv -ggdb3 -o stdin-example stdin-e
echo 10 | ./stdin-example 2>&1
echo 20 > 20.txt
echo 30 > 30.txt
./stdin-example 2>&1 < 20.txt
./stdin-example < 30.txt 2>&1
```

```
ERR: From stdin 10
From stdin 10
ERR: From stdin 20
From stdin 20
ERR: From stdin 30
From stdin 30
```

You’ve mostly seen this before except the ‘`fprintf(stderr,...)`’ part.

1.5.2 Files

Files can be addressed as streams as well. But we have to open and close them. So we can treat files exactly like `stdin` and `stdout` but with a few changes.

1. We need a file handle (like `stdin`, `stdout`, or `stderr`). This handle is for the OS to know which file the process is talking about.
 - `fopen`
2. We need to decide if we are reading write or both and we need to open a file to produce a file handle. Or use an existing one.
 - `fprintf`
 - `fputs`
3. We need to write to it using `write` and `f*` operations.
 - `fprintf`
 - `fputs`
4. We need to read from it using `read` and `f*` operations.
 - `fgets`
 - `fgetc`
5. We need to close the file after we’re done. `fclose`.

1. open and close

To open a file we use `fopen`. To close it we `fclose`. Don't use `open` and `close` because that's not portable. That's for the OS.

```
FILE *fopen(const char *pathname, const char *mode);
int fclose(FILE *stream);
```

The mode is a string

- "r" - read
- "w" - write (erase file)
- "a" - append (add to end of file)
- "r+" - read and write
- "w+" - write and read (erase file)
- "a+" - append and read

```
FILE * f_cb = fopen("coolbears.txt", "w"); // open coolbears.txt for writing
int fclose(f_cb); // close coolbears.txt
```

If you don't close a file you can lose bytes you wrote to it because they didn't get flushed to disk. This is important because people might kill your program your you might reboot or shutdown the computer. If you want to ensure data is written try to engage in `flush`. Sometimes no data will appear until you flush or close the file. Keep those pipes clean.

```
#include <stdio.h>
#define SIZE 1024
int main() {
    char buffer[SIZE] = {'\0'};
    // open coolbears.txt for writing
    FILE * f_cb = fopen("coolbears.txt", "w");
    fputs("Polar bears", f_cb);
    fclose(f_cb);
    FILE * f_cbr = fopen("coolbears.txt", "r");
    fgets(buffer, SIZE, f_cbr);
    printf("%s\n",buffer);
    fclose(f_cbr);
}
```

Polar bears

OK but what if we don't close it?

```
#include <stdio.h>
#define SIZE 1024
int main() {
    char buffer[SIZE] = {'\0'};
    // open coolbears.txt for writing
    FILE * f_cb = fopen("coolbears.txt", "w");
    fputs("Polar bears", f_cb);
    FILE * f_cbr = fopen("coolbears.txt", "r");
    fgets(buffer, SIZE, f_cbr);
    printf("This is the buffer before close: %s\n",buffer);
    fclose(f_cbr);
    fclose(f_cb);
    f_cbr = fopen("coolbears.txt", "r");
    fgets(buffer, SIZE, f_cbr);
    printf("This is the buffer after close: %s\n",buffer);
    fclose(f_cbr);
    printf("Close your buffers!");
}
```

```
This is the buffer before close:
This is the buffer after close: Polar bears
Close your buffers!
```

```
cat coolbears.txt
```

Polar bears

Now let's see what flush does for us!

```
#include <stdio.h>
#define SIZE 1024
int main() {
    char buffer[SIZE] = {'\0'};
    // open coolbears.txt for writing
    FILE * f_cb = fopen("coolbears.txt", "w");
```

```

    fputs("Polar bears", f_cb);
    fflush(f_cb); // WE'RE FLUSHING!
    FILE * f_cbr = fopen("coolbears.txt", "r");
    fgets(buffer, SIZE, f_cbr);
    printf("This is the buffer before close but after flush: %s\n",buffer);
    fclose(f_cbr);
    fclose(f_cb);
    f_cbr = fopen("coolbears.txt", "r");
    fgets(buffer, SIZE, f_cbr);
    printf("This is the buffer after close: %s\n",buffer);
    fclose(f_cbr);
    printf("Close your buffers! Keep your pipes clean");
}

```

This is the buffer before close but after flush: Polar bears
 This is the buffer after close: Polar bears
 Close your buffers! Keep your pipes clean

2. writing

(a) fprintf

```

#include <stdio.h>
#include <stdlib.h>
#define SIZE 1024
int main() {
    srand(time(NULL));
    char buffer[SIZE] = {'\0'};
    // open coolbears.txt for writing
    FILE * f_cb = fopen("fprintf.txt", "w");
    // It's just like printf!
    fprintf(f_cb,"A random number %d\n", rand());
    fclose(f_cb);
    FILE * f_cbr = fopen("fprintf.txt", "r");
    fgets(buffer, SIZE, f_cbr);
    printf("%s\n",buffer);
    fclose(f_cbr);
}

```

A random number 440052301

3. reading

(a) fscanf

```
#include <stdio.h>
#include <stdlib.h>
#define SIZE 1024
#define CHECK(x) ((x)==1)?1:(abort(),0);
int main() {
    srand(time(NULL));
    char buffer[SIZE] = {'\0'};
    // open coolbears.txt for writing
    FILE * f_cb = fopen("fprintf.txt", "w");
    // It's just like printf!
    fprintf(f_cb, "A random number %d\n", rand());
    fclose(f_cb);
    FILE * f_cbr = fopen("fprintf.txt", "r");
    for (int i = 0 ; i < 3; i++) {
        CHECK(fscanf(f_cbr, "%s", buffer));
        printf("%s\n", buffer);
    }
    int input=0;
    CHECK(fscanf(f_cbr, "%d",&input));
    printf("%d\n", input);
    fclose(f_cbr);
    return 0;
}

A
random
number
440052301
```

(b) fgets

fgets gets a little complicated because you have to test for EOF. You can check for an null response and use feof, but probably you have to do both.

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
```

```

#define SIZE 1024
#define N 8
int main() {
    char buffer[SIZE] = {'\0'};
    srand(time(NULL));
    // open coolbears.txt for writing
    FILE * f_cb = fopen("fgets.txt", "w");
    // It's just like printf!
    const int totalLines = 1 + (rand() % N);
    for (int i = 0 ; i < totalLines; i++) {
        fprintf(f_cb, "A random number %d\n", rand());
    }
    fclose(f_cb);
    FILE * f_cbr = fopen("fgets.txt", "r");
    while(!feof(f_cbr)) {
        if (fgets(buffer, SIZE, f_cbr)) {
            printf("fgets.txt: %s", buffer);
        }
    }
    fclose(f_cbr);
    return 0;
}

fgets.txt: A random number 502032220
fgets.txt: A random number 701620934
fgets.txt: A random number 1131098919
fgets.txt: A random number 176019583
fgets.txt: A random number 1431369748
fgets.txt: A random number 118249612

```

4. flushing

If you want to ensure something gets to a file or a term you should flush. Typically I/O is BUFFERED. Buffered means it is flushed once a certain threshold is met, typically size but sometimes time (depending on the system). Buffered will increase latency to print something but will often improve overall bandwidth to disk.

```
fflush(FILE * stream); // will flush your stream
```

Flush when you need to.

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <unistd.h>
#define SIZE 5
int main() {
    srand(time(NULL));
    FILE * file = fopen("fflush.txt", "w");
    for (int i = 0 ; i < SIZE; i++) {
        fprintf(file, "%d\n", rand());
        fflush(file); // WE'RE FLUSHING!
        sleep(1);
    }
    fclose(file);
}

```

```

#include <time.h>
#include <stdio.h>
#include <unistd.h>
#define SIZE 20
#define BUFF 1024
int main() {
    char buffer[BUFF];
    FILE * file = fopen("fflush.txt", "r");
    while(!feof(file)) {
        if (!fgets(buffer,BUFF,file)) {
            break;
        }
        printf("%s", buffer);
        sleep(1);
    }
    fclose(file);
}

```

```

gcc -std=c99 -pedantic -Wall -Wextra -ftrapv -ggdb3 -o fflushrandr fflushrandr.c
gcc -std=c99 -pedantic -Wall -Wextra -ftrapv -ggdb3 -o fflushreader fflushreader.c
echo This will take 7 seconds && \
( ./fflushrandr & sleep 2; ./fflushreader)

```

This will take 7 seconds

```
440052301
502032220
701620934
1131098919
176019583
```

5. Binary Files

From `stdio.h`:

```
size_t fread(void *ptr, size_t size, size_t nmemb, FILE *stream);

size_t fwrite(const void *ptr, size_t size, size_t nmemb,
               FILE *stream);
```

`fread` and `fwrite` will write memory to a stream and back again. Any pointer can be used, the bytes in memory will be serialized in and out. It will not be compiler and architecture portable so carefully craft your structs before you write them out. Use explicit padding. For 64-bit and 32-bit compatibility pad to modulus 8 bytes.

```
#include <assert.h>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#define SIZE 5
struct demo {
    int i;
    float f;
    double d;
    char c;
};
int main() {
    srand(time(NULL));
    FILE * file = fopen("binary.bin", "w");
    int realSize = 1 + (rand() % SIZE);
    assert(1==fwrite(&realSize, sizeof(int), 1, file));
    for (int i = 0 ; i < realSize; i++) {
        struct demo randd;
        randd.i = rand();
```

```

        randd.f = rand() / 2.0F;
        randd.d = 1.0 / (rand()+1.0);
        randd.c = 'X';
        printf("Writing %d\n", randd.i);
        printf("\tWriting %g\n", randd.f);
        printf("\tWriting %g\n", randd.d);
        printf("\tWriting %c\n", randd.c);
        assert(1==fwrite(&randd, sizeof(randd), 1, file));
    }
    fclose(file);
}

#include <assert.h>
#include <time.h>
#include <stdio.h>
#define BUFF 1024
struct demo {
    int i;
    float f;
    double d;
    char c;
};
int main() {
    char buffer[BUFF];
    FILE * file = fopen("binary.bin", "r");
    int size=0;
    assert(1==fread(&size, sizeof(size), 1, file));

    while(!feof(file)) {
        struct demo readDemo;
        if (1!=fread(&readDemo, sizeof(readDemo), 1, file)) {
            break;
        }
        printf("Reading %d\n", readDemo.i);
        printf("\tReading %f\n", readDemo.f);
        printf("\tReading %g\n", readDemo.d);
        printf("\tReading %c\n", readDemo.c);
    }
    fclose(file);
}

```


When you write out structs, watch for padding. Look for the letter X.
Count the number of bytes after the last X.

```
hexdump -C binary.bin
```

```
00000000  04 00 00 00 a0 89 b3 49 61 d2 42 4d 66 42 7d d0 |.....Ia.BMfB}.|
00000010  62 14 60 3e 58 7b 14 70 05 56 00 00 e8 da 0c 70 |b.'>X{.p.V.....p|
00000020  47 da 1e 4e fb 8c cf 16 64 0b 00 3e 58 7b 14 70 |G..N....d..>X{.p|
00000030  05 56 00 00 35 80 92 1b 36 55 52 4e 72 e9 b3 c3 |.V..5...6URNr...|
00000040  b6 f2 0a 3e 58 7b 14 70 05 56 00 00 ad db 72 2a |...>X{.p.V....r*|
00000050  7b 64 fe 4d e8 a4 a0 06 a7 e1 0e 3e 58 7b 14 70 |{d.M.....>X{.p|
00000060  05 56 00 00                                     |.V...|
00000064
```

1.5.3 Command line arguments

To get arguments from the commandline you can add the parameters:

- ‘int argc’ – number of commandline arguments
- ‘char ** argv’ – array of strings of commandline arguments

```
#include <stdio.h>
```

```
int main(int argc, char ** argv) {
    for (int i = 0 ; i < argc; i++) {
        printf("arg %d: %s\t", i, argv[i]);
    }
    puts("");
}
```

```
arg 0: /tmp/babel-7888jxN/C-bin-7888GIO
```

```
gcc -std=c99 -pedantic -Wall -Wextra -ftrapv -ggdb3 -o argv argv.c && \
./argv && \
./argv 1 && \
./argv 1 2 && \
./argv 1 2 3 && \
cp -f argv argv-new && \
./argv-new 1 2 3
```

```

arg 0: ./argv
arg 0: ./argv arg 1: 1
arg 0: ./argv arg 1: 1 arg 2: 2
arg 0: ./argv arg 1: 1 arg 2: 2 arg 3: 3
arg 0: ./argv-new arg 1: 1 arg 2: 2 arg 3: 3

```

So atoi is your friend :)

```

#include <stdio.h>
#include <stdlib.h>

int main(int argc, char ** argv) {
    if (argc != 2) { exit(1); }
    int n = atoi(argv[1]);
    for (int i = 0 ; i < n; i++) {
        printf("%d\t", rand());
    }
    printf("\n");
}

gcc -std=c99 -pedantic -Wall -Wextra -ftrapv -ggdb3 \
    -o argvrand argvrand.c && \
(./argvrand || echo not enough args: $?) && \
./argvrand 1 && \
./argvrand 2 && \
./argvrand 3 && \
./argvrand 0

not enough args: 1
1804289383
1804289383 846930886
1804289383 846930886 1681692777

```

1.5.4 mmap()

mmap is neat, it maps memory from a file or even a process to pointers in memory. We do this with shared libraries too. So you can read and write to a file just by writing to memory. The OS deals with it very efficiently you just have to be very size aware. mmaping files is not good for streams and stream processing, it gets complicated. It is good for fixed sized structs.

```

#include <assert.h>
#include <sys/mman.h>
#include <sys/types.h>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#define SIZE 2
struct demo {
    int i;
    float f;
    double d;
    char c;
    // char cc[7]; // you can make padding explicit
};
int main() {
    srand(time(NULL));
    FILE * file = fopen("binary.bin", "r");
    int fd = fileno(file);
    int rsize = 0;
    assert(1==fread(&rsize, sizeof(rsize), 1, file));
    const size_t size = sizeof(struct demo) * rsize;
    printf("N %d struct demos are in binary.bin\n", rsize);
    printf("mmapping %u bytes of memory from the file\n", size);
    int * mapped = mmap(0,
        size,
        PROT_READ, // | PROT_WRITE,
        MAP_SHARED,
        fd,
        0);
    if (mapped == MAP_FAILED) {
        perror("mmap");
        exit(1);
    }
    // read 4 bytes from the head
    assert(rsize == (int)*mapped);
    // really abusive but we're 1 int after the start eh
    struct demo * demos = (struct demo *) (mapped+1);
    // OK now look for the read ?
    for (int i = 0 ; i < rsize; i++) {
        struct demo randd = demos[i];
    }
}

```

```

        printf("Reading %d\n", randd.i);
        printf("\tReading %g\n", randd.f);
        printf("\tReading %g\n", randd.d);
        printf("\tReading %c\n", randd.c);
    }
    munmap(demos, size);
    fclose(file);
}

```

If you want to see some bad code that's small and uses mmap checkout:

<https://github.com/abramhindle/a-simple-pseudo-bayesian-spam-filter/blob/master/filter.c>

1.6 References

KN King, C Programming, Chapter 28, 2nd Edition

Hazel Cambell's thorough notes on Stream I/O: <https://docs.google.com/document/d/1b48EzfP03JYEft42wCajU5kv76oVbTxEXa2J00q17ag/edit>